


Text Mining


Natural Language Processing &
Semantic Computing



Prof. Dr. Siegfried Handschuh

Text Mining

Basic Text Processing



Prof. Dr. Siegfried Handschuh

Basic Text Processing

- Regular Expression,
- Regular Expressions in Practical NLP
- Word Tokenisation
- Word Normalisation and Stemming,
- Sentence Segmentation

Exercises

- Due to huge demand, we will have one big Exercises session!
- Bernhard will explain
- Individual Feedback via Piazza Learning System (you will get an introduction to that)

Regular Expressions

Regular expressions

- A formal language for specifying text strings
- How can we search for any of these?
 - woodchuck
 - woodchucks
 - Woodchuck
 - Woodchucks



Regular Expressions: Disjunctions

Letters inside square brackets []

Pattern	Matches
<i>[wW]oodchuck</i>	<i>Woodchuck, woodchuck</i>
<i>[1234567890]</i>	<i>Any digit</i>

Ranges ***[A-Z]***

Pattern	Matches	
<i>[A-Z]</i>	<i>An upper case letter</i>	<i><u>D</u>renched Blossoms</i>
<i>[a-z]</i>	<i>A lower case letter</i>	<i><u>m</u>y beans were impatient</i>
<i>[0-9]</i>	<i>A single digit</i>	<i>Chapter <u>1</u>: Down the Rabbit Hole</i>

Regular Expressions: Negation in Disjunction

Negations [^Ss]

Caret means negation only when first in []

Pattern	Matches	
[^A-Z]	<i>Not an upper case letter</i>	O <u>y</u> fn pripetchik
[^Ss]	<i>Neither 'S' nor 's'</i>	<u>I</u> have no exquisite reason"
[^e^]	<i>Neither e nor ^</i>	<u>L</u> ook here
a \ ^b	<i>The pattern a caret b</i>	Look up <u>a^b</u> now

Regular Expressions: More Disjunction

- Woodchuck is another name for groundhog!
- The pipe | for disjunction


Pattern	Matches
<i>groundhog/woodchuck</i>	
<i>yours/mine</i>	<i>yours</i> <i>mine</i>
<i>a/b/c</i>	<i>= [abc]</i>
<i>[gG]roundhog/ [Ww]oodchuck</i>	



Regular Expressions:

? * + .

Pattern	Matches	
<i>colou?r</i>	<i>Optional previous char</i>	<u>color</u> <u>colour</u>
<i>oo*h!</i>	<i>0 or more of previous char</i>	<u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>
<i>o+h!</i>	<i>1 or more of previous char</i>	<u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>
<i>baa+</i>		<u>baa</u> <u>baaa</u> <u>baaaa</u> <u>baaaaa</u>
<i>beg.n</i>		<u>begin</u> <u>begun</u> <u>begun</u> <u>beg3n</u>



Stephen C Kleene

Kleene *, Kleene +


Regular Expressions: Anchors

^ **\$**

Pattern	Matches
^[A-Z]	<u>P</u> alo Alto
^[^A-Za-z]	<u>1</u> <u>"Hello"</u>
\. \$	The end <u>.</u>
. \$	The end <u>?</u> The end <u>!</u>

Excursus: Dr Seuss

Seuss's **cat** in the hat



Example: Cat in the Hat

We looked!
Then we saw him step in on the mat.
We looked!
And we saw him!
The Cat in the Hat!

The other one there, the blithe one.

Find me all instances of the word “the” in a text.

the

Misses capitalised example

[tT]he

Incorrectly returns other or blithe

[^a-zA-Z][tT]he[^a-zA-Z]

Errors

The process we just went through was based on **fixing two kinds of errors**

- Matching strings that we should not have matched (**there, then, other**)
 - **False positives (Type I)**
- Not matching things that we should have matched (The)
 - **False negatives (Type II)**

Errors cont.

- In NLP we are always dealing with these kinds of errors.
- Reducing the error rate for an application often involves two antagonistic efforts:
 - Increasing accuracy or precision
(minimising false positives)
 - Increasing coverage or recall
(minimising false negatives).

Summary

- Regular expressions play a surprisingly large role
 - Sophisticated sequences of regular expressions are often the first model for any text processing text
- For many hard tasks, we use machine learning classifiers
 - But regular expressions are used as features in the classifiers
 - Can be very useful in capturing generalisations

Assignment

- write regular expressions that extract **phone numbers** and regular expressions that extract **email addresses**.
- i.e. peter.mueller@uni-passau.de
- + also variations, such as:
peter dot mueller at uni-passau dot de, etc.
- + java script encodings, such as used by the university of passau

Word tokenisation

Text Normalisation

- Every NLP task needs to do text normalisation:
 1. Segmenting/tokenising words in running text
 2. Normalising word formats
 3. Segmenting sentences in running text

How many words?

- I do uh main- mainly business data processing
 - Fragments, filled pauses
- Seuss's **cat** in the hat is different from other **cats**!
 - **Lemma**: same stem, part of speech, rough word sense
 - **cat** and **cats** = same lemma
 - **Wordform**: the full inflected surface form
 - **cat** and **cats** = different wordforms

How many words?

they lay back on the San Francisco grass and looked at the stars and their

- **Type**: an element of the vocabulary.
- **Token**: an instance of that type in running text.
- How many?
 - 15 tokens (or 14)
 - 13 types (or 12) (or 11?)

QUIZ

1. How many wordform types and tokens are in this sentence:

I want to cook a dish that I never cooked before

- ☐ A 10 types 11 tokens
- ☐ B 11 types 10 tokens
- ☐ C 11 types 11 tokens
- ☐ D 11 types 10 tokens

How many words?

N = number of tokens

V = vocabulary = set of types

$|V|$ is the size of the vocabulary

	Tokens = N	Types = $ V $
<i>Switchboard phone</i>	<i>2.4 million</i>	<i>20,000</i>
<i>Shakespeare</i>	<i>884,000</i>	<i>31,000</i>
<i>Google N-grams</i>	<i>1 trillion</i>	<i>13 million</i>

Simple Tokenization in UNIX

(Inspired by Ken Church's UNIX for Poets.)

Given a text file, output the word tokens and their frequencies

```
tr -sc 'A-Za-z' '\n' < shakes.txt
```

Change all non-alpha to newlines

```
| sort
```

Sort in alphabetical order

```
| uniq -c
```

Merge and count each type

```
1
17 A
1 ADVENTURER
1 ALL
1 AND
1 Above
1 Accuse
1 Admit
1 Adonis
```


The first step: tokenising

```
tr -sc 'A-Za-z' '\n' < shakes.txt | head
```

```
THE  
SONNETS  
by  
William  
Shakespeare  
From  
fairest  
creatures  
We  
...
```

The second step: sorting

```
tr -sc 'A-Za-z' '\n' < shakes.txt | sort | head
```

A

A

A

A

A

A

A

A

A

...

More Counting

Merging upper and lower case

```
tr 'A-Z' 'a-z' < shakes.txt | tr -sc 'A-Za-z' '\n' | sort
```

Sorting the counts

```
tr 'A-Z' 'a-z' < shakes.txt | tr -sc 'A-Za-z' '\n' | sort |  
uniq -c | sort -n -r
```

```
322 that  
266 thy  
235 thou  
200 s  
195 love  
181 with  
172 for  
170 is  
168 not  
166 d
```

What happened here?

QUIZ

2. We saw that the words 'd' and 's' appear frequently in the word tokenization results. Which of following reasons might explain this?

(A) The words "s" and "d" are archaic words that were more popular during the time that Shakespeare was writing.

(B)

The words "s" and "d" were not very common during the time that Shakespeare was writing, but they are among the many words which Shakespeare appears to have made up himself.

(C)

The tokenizer splits on all non-alphabetical characters, which includes the apostrophe. This mean that words like "Ophelia's" and "'dimm'd" will be tokenized to "Ophelia s" and "dimm d".

(D)

The Tokenizer splits words into smaller subparts named morphemes. "s" and "d" are two very common such morphemes because they tend to occur at the end of words.

Shakespeare

Will be a tatter'd weed,
Then being ask'd where
For where is she so fair whose unear'd womb
But if thou live, remember'd not to be,

Issues in Tokenisation


- Finland's capital → Finland Finlands Finland's ?
- what're, I'm, isn't → What are, I am, is not
- Hewlett-Packard → Hewlett Packard ?
- state-of-the-art → state of the art ?
- Lowercase → lower-case lowercase lower case ?
- San Francisco → one token or two?
- m.p.h., PhD. → ??

Tokenisation: language issues

- French
 - **L'ensemble** → one token or two?
 - **L ? L' ? Le ?**
 - Want **l'ensemble** to match with **un ensemble**
- German noun compounds are not segmented
 - **Versicherungsangestellter** =
insurance company employee
 - **Donaudampfschiffahrtskapitän** =
danube steamship captain'
 - German information retrieval needs compound splitter

Tokenisation: language issues

- Chinese and Japanese no spaces between words:
 - 莎拉波娃现在居住在美国东南部的佛罗里达。
 - 莎拉波娃 现在 居住 在 美国 东南部 的 佛罗里达
 - Sharapova now lives in US southeastern Florida
- Further complicated in Japanese, with multiple alphabets intermingled
 - Dates/amounts in multiple formats



End-user can express query entirely in hiragana!

Word Tokenisation in Chinese

- Also called **Word Segmentation**
- Chinese words are composed of characters
 - Characters are generally 1 syllable and 1 morpheme.
 - Average word is 2.4 characters long.
- Standard baseline segmentation algorithm:
 - Maximum Matching (also called Greedy)

Maximum Matching Word Segmentation Algorithm

- Given a wordlist of Chinese, and a string.
 1. Start a pointer at the beginning of the string
 2. Find the longest word in dictionary that matches the string starting at pointer
 3. Move the pointer over the word in string
 4. Go to 2

Max-match segmentation illustration

- Thecatinthehat the cat in the hat
- Thetabledownthere the table down there
 theta bled own there
- Doesn't generally work in English!
- But works astonishingly well in Chinese
 - 莎拉波娃现在居住在美国东南部的佛罗里达。
 - 莎拉波娃 现在 居住 在 美国 东南部 的 佛罗里达
- Modern probabilistic segmentation algorithms even better

Word Normalisation and Stemming

Normalisation

- Need to “normalise” terms
 - Information Retrieval: indexed text & query terms must have same form.
 - We want to match ***U.S.A.*** and ***USA***
- We implicitly define equivalence classes of terms
 - e.g., deleting periods in a term
- Alternative: asymmetric expansion:
 - Enter: ***window*** Search: ***window, windows***
 - Enter: ***windows*** Search: ***Windows, windows, window***
 - Enter: ***Windows*** Search: ***Windows***
- Potentially more powerful, but less efficient

Case folding

- Applications like IR: reduce all letters to lower case
 - Since users tend to use lower case
 - Possible exception: upper case in mid-sentence?
 - e.g., *General Motors*
 - *Fed* vs. *fed*
 - *SAIL* vs. *sail*
- For sentiment analysis, MT, Information extraction
 - Case is helpful (*US* versus *us* is important)

Lemmatisation

- Reduce inflections or variant forms to base form
 - *am, are, is* → *be*
 - *car, cars, car's, cars'* → *car*
- *the boy's cars are different colours* → *the boy car be different colour*
- Lemmatisation: have to find correct dictionary headword form
- Machine translation
 - Spanish **quiero** ('I want'), **quieres** ('you want') same lemma as **querer** 'want'

Morphology

- **Morphemes:**
 - The small meaningful units that make up words
 - **Stems:** The core meaning-bearing units
 - **Affixes:** Bits and pieces that adhere to stems
 - Often with grammatical functions

Stemming

- Reduce terms to their stems in information retrieval
- *Stemming* is crude chopping of affixes
 - language dependent
 - e.g., ***automate(s), automatic, automation*** all reduced to ***automat***.

*for example compressed
and compression are both
accepted as equivalent to
compress.*



for exampl compress and
compress ar both accept
as equal to compress

Porter's algorithm: The most common English stemmer

Step 1a

sses	→ ss	caresses	→ caress
ies	→ i	ponies	→ poni
ss	→ ss	caress	→ caress
s	→ ∅	cats	→ cat

Step 1b

(*v*)ing	→ ∅	walking	→ walk
		sing	→ sing
(*v*)ed	→ ∅	plastered	→ plaster
...			

Step 2 (for long stems)

ational	→ ate	relational	→ relate
izer	→ ize	digitizer	→ digitize
ator	→ ate	operator	→ operate
...			

Step 3 (for longer stems)

al	→ ∅	revival	→ reviv
able	→ ∅	adjustable	→ adjust
ate	→ ∅	activate	→ activ
...			

QUIZ

3. Given the description you saw on earlier slides, the Porter stemmer would stem the word 'aching' as
- ☒ A aching
 - ☐ B ach
 - ☐ C ache
 - ☐ D aches

Why only strip -ing if there is a vowel?

(*v*)ing → ∅ walking → walk
 sing → sing

Why only strip -ing if there is a vowel?

(*v*)ing → ∅ walking → walk
 sing → sing

```
tr -sc 'A-Za-z' '\n' < shakes.txt | grep 'ing$' | sort | uniq -c | sort -nr
```

```

1312 King
548 being
541 nothing
152 something
388 king
145 coming
375 bring
130 morning
358 thing
122 having
307 ring
120 living
152 something
117 loving
145 coming
116 Being
130 morning
102 going
548 being
541 nothing
152 something
145 coming
130 morning
122 having
120 living
117 loving
116 Being
102 going

```

tr -sc 'A-Za-z' '\n' < shakes.txt | grep '[aeiou]*ing\$' | sort | uniq -c | sort -nr

Viewing morphology in a corpus

Why only strip –ing if there is a vowel?

(*v*)ing → ∅ walking → walk
 sing → sing

```
tr -sc 'A-Za-z' '\n' < shakes.txt | grep 'ing$' | sort | uniq -c | sort -nr
```

1312 King	548 being
548 being	541 nothing
541 nothing	152 something
388 king	145 coming
375 bring	130 morning
358 thing	122 having
307 ring	120 living
152 something	117 loving
145 coming	116 Being
130 morning	102 going

```
!tr -sc 'A-Za-z' '\n' < shakes.txt | grep '[aeiou].*ing$' | sort | uniq  
-c | sort -nr
```

Dealing with complex morphology is sometimes necessary

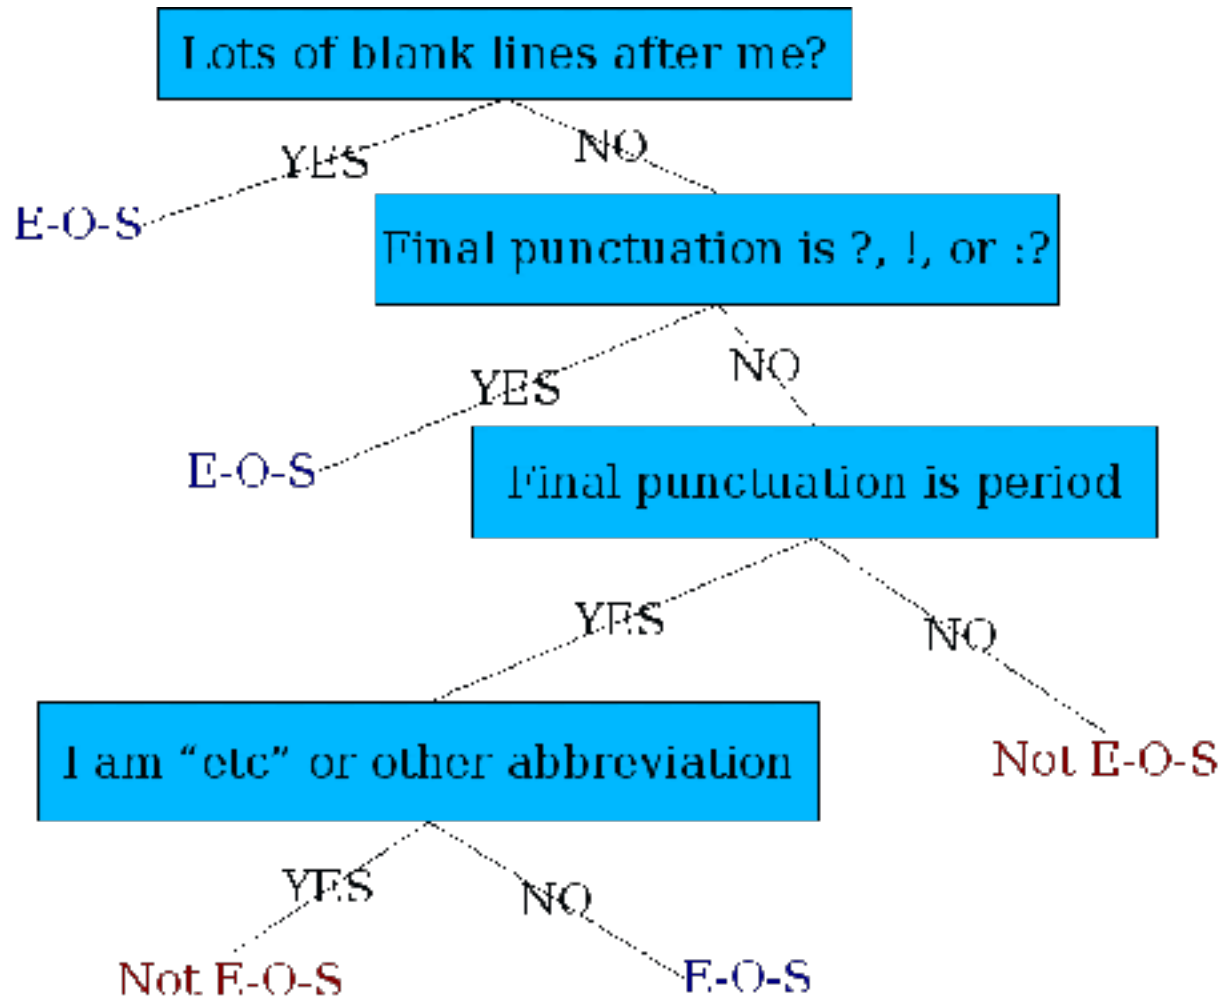
- Some languages requires complex morpheme segmentation
 - Turkish
 - **Uygarlastiramadiklarimizdanmissinizcasina**
 - `(behaving) as if you are among those whom we could not civilised'
 - **Uygar** `civilized' + **las** `become'
 - + **tir** `cause' + **ama** `not able'
 - + **dik** `past' + **lar** `plural'
 - + **imiz** `plpl' + **dan** `abl'
 - + **mis** `past' + **siniz** `2pl' + **casina** `as if'

Sentence Segmentation

Sentence Segmentation

- !, ? are relatively unambiguous
- Period “.” is quite ambiguous
 - Sentence boundary
 - Abbreviations like Inc. or Dr.
 - Numbers like .02% or 4.3
- Build a binary classifier
 - Looks at a “.”
 - Decides EndOfSentence/NotEndOfSentence
 - Classifiers: hand-written rules, regular expressions, or machine-learning

Determining if a word is end-of-sentence: a Decision Tree



More sophisticated decision tree features

- Case of word with “.”: Upper, Lower, Cap, Number
- Case of word after “.”: Upper, Lower, Cap, Number
- Numeric features
 - Length of word with “.”
 - Probability(word with “.” occurs at end-of-s)
 - Probability(word after “.” occurs at beginning-of-s)

QUIZ

A period (".") occurs at the end of four words that each have one of the following four wordshapes. Which of the four periods (all else being equal) is more likely than the other to be a sentence boundary?

- (A) Upper
- (B) Lower
- (C) Cap
- (D) Number

QUIZ

- A: "... Word."
- B: "... word." most likely to be EOS (End Of Sentence)
- C "...WORD."
- D "... 9."

Implementing Decision Trees

- A decision tree is just an if-then-else statement
- The interesting research is choosing the features
- Setting up the structure is often too hard to do by hand
 - Hand-building only possible for very simple features, domains
 - For numeric features, it's too hard to pick each threshold
 - Instead, structure usually learned by machine learning from a training corpus

Decision Trees and other classifiers

- We can think of the questions in a decision tree
- As features that could be exploited by any kind of classifier
 - Logistic regression
 - SVM
 - Neural Nets
 - etc.

Basic Text Processing

- Regular Expression,
- Regular Expressions in Practical NLP
- Word Tokenisation
- Word Normalisation and Stemming,
- Sentence Segmentation